



MODULE 9: INTRODUCTION TO DAE

BA713 - Machine Learning & AI

CONTENTS



GENERATIVE MODELING



INTRODUCTION TO DAE



ANOMALY DETECTION

SUPERVISED VERSUS UNSUPERVISED LEARNING-Recap

Supervised Learning

- **Data:** features + labels
- **Objective:** Learn to match the features with labels
- **Example:** Classification, regression, Object detection, etc.

Unsupervised Learning

- **Data:** only features (no labels)
- **Objective:** Learn underlying data structure and relationships
- **Example:** Clustering, Feature selection, dimensionality reduction, etc.

GENERATIVE VERSUS DISCRIMINATIVE MODELS

Generative models:

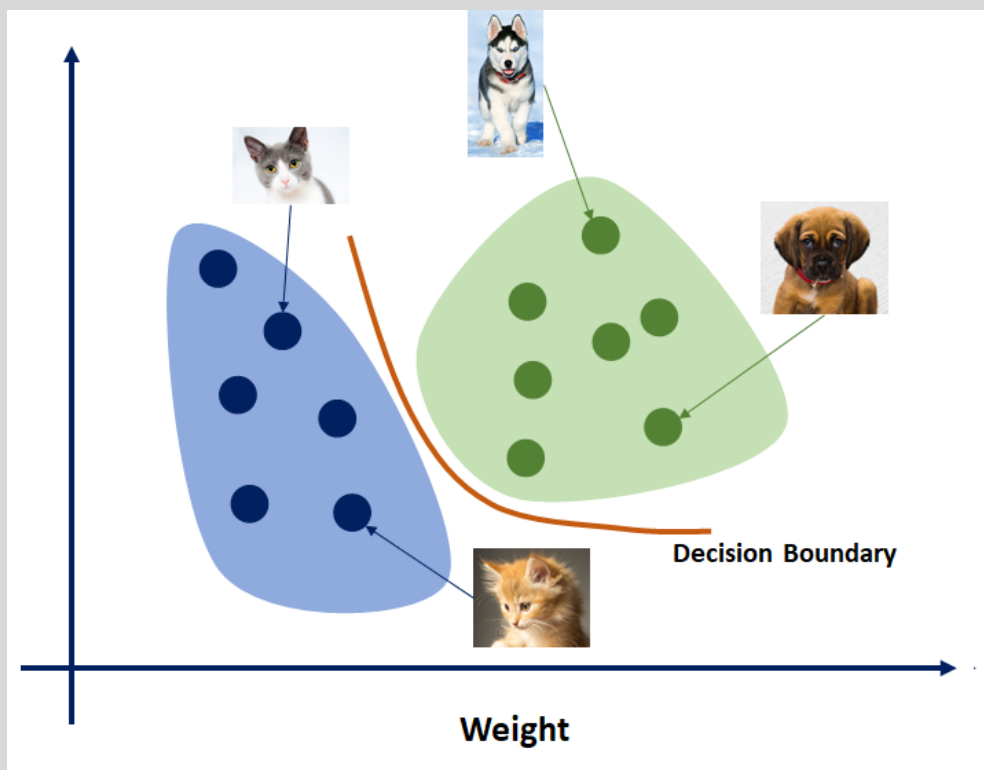
- can generate new data instances.
- learn the data distribution using **unsupervised learning**
- aim at learning the true data distribution of the training observations to generate new observations with some variations
- Exact distribution → not always possible
- given a set of data instances X and a set of labels Y
- capture the joint probability $p(X, Y)$, or just $p(X)$ if no labels are available
- For example, model that can predict the next word in a sequence
- **Example models** → GAN and its variants, DAE and its variants

Discriminative models:

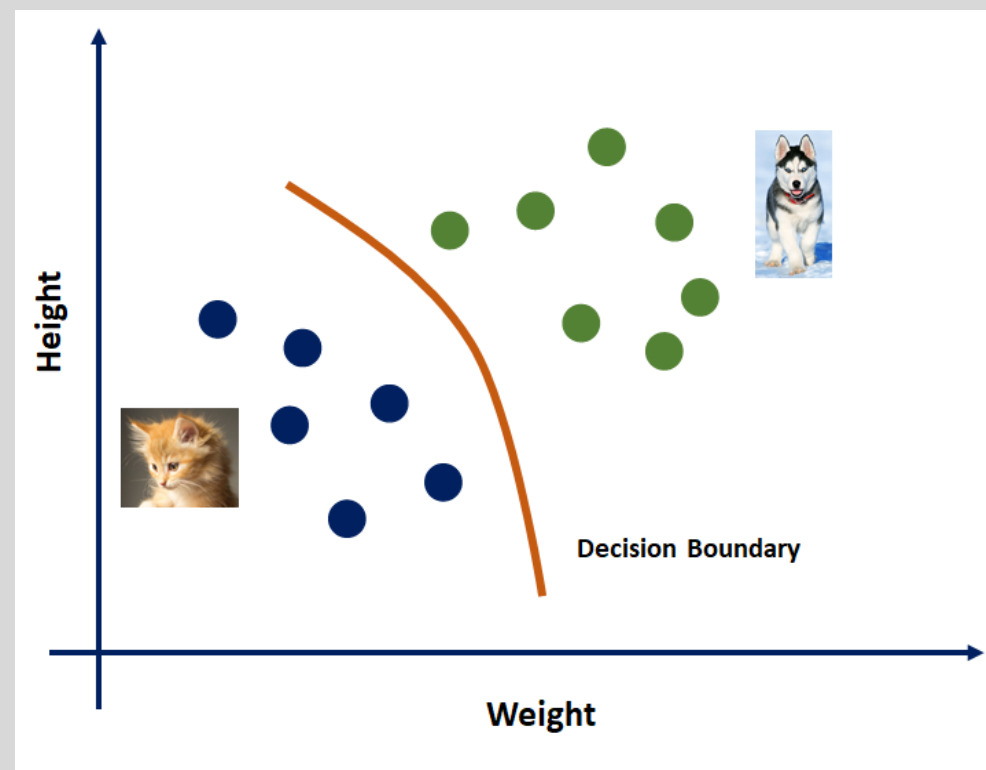
- discriminate between different kinds of data instances
- **supervised learning**: e.g., classification
- capture the conditional probability $p(Y | X)$
- For example, how likely a particular label can apply to a specific observation/instance

GENERATIVE VERSUS DISCRIMINATIVE MODELS

Generative Model

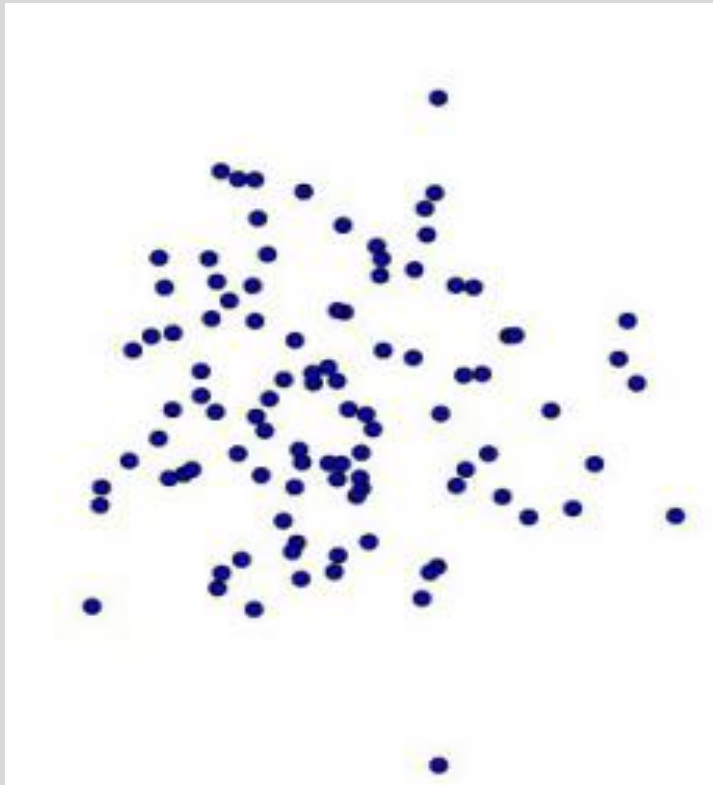


Discriminative Model

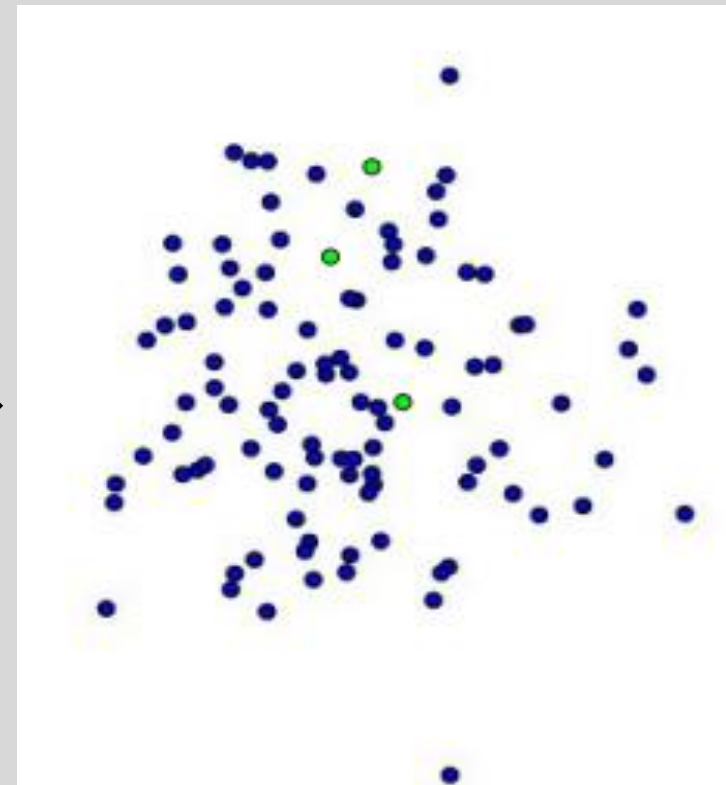
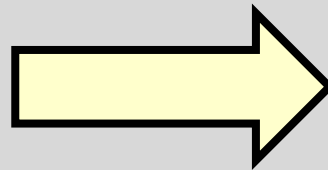


GENERATIVE MODELING

- Generate samples with the same probabilistic distribution as input samples



**Input data samples in
blue color**



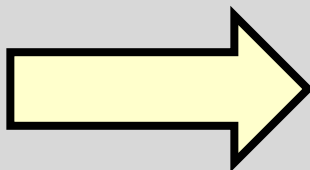
**Generated samples in
green color**

NEED FOR GENERATIVE MODELING

- Learn underlying features in the dataset without the need for labels
- Representative datasets
- More learning samples
- Solve the class imbalance problem

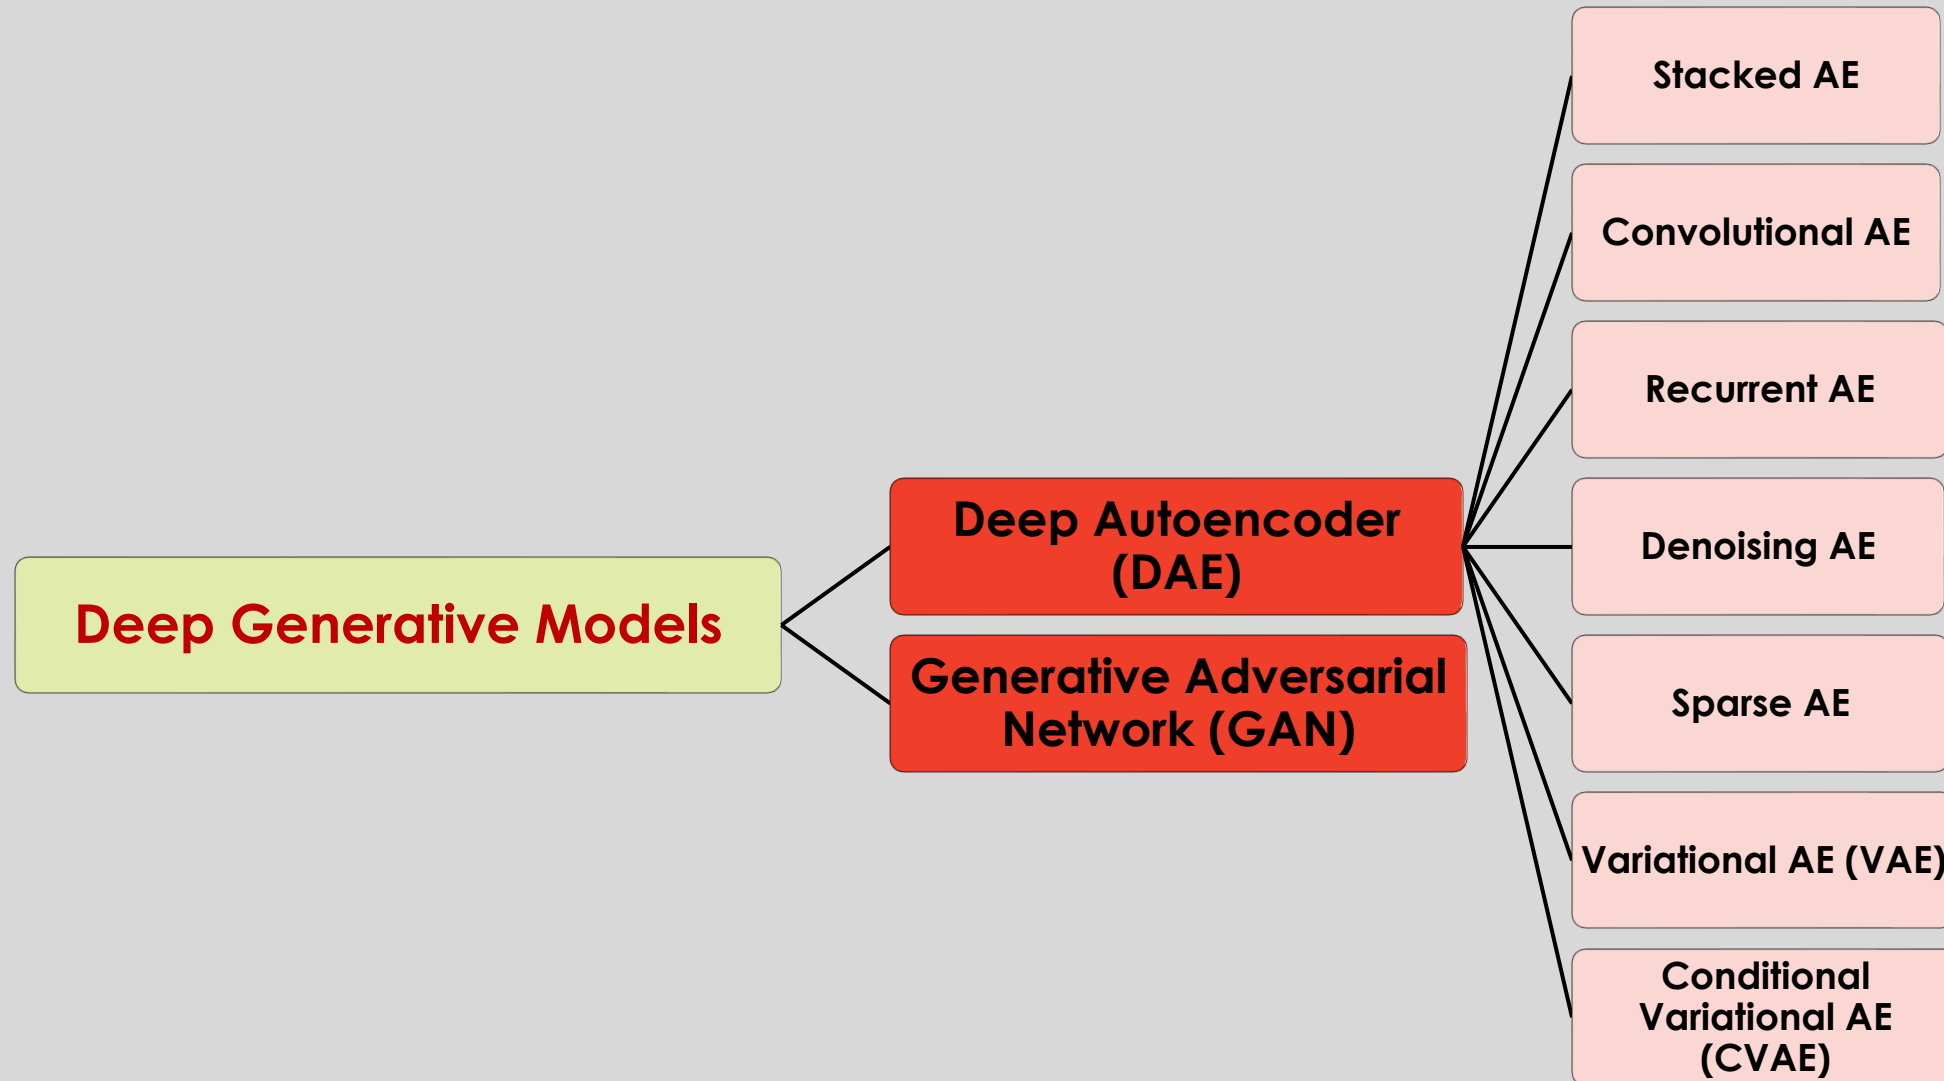


Homogenous skin color and pose



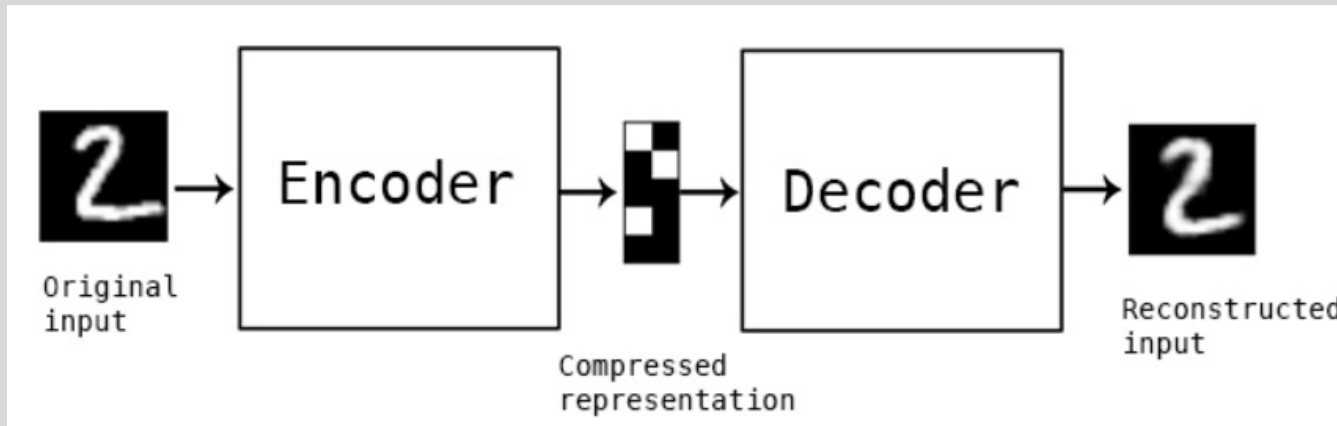
Diverse skin color, pose, and illumination

DEEP GENERATIVE MODELS CLASSIFICATION



AUTOENCODER MODEL

- data compression algorithm
- Implemented with neural networks
- **Self-supervised learning** → Supervised Learning with automatically generated labels
- **Data-specific** → compress data similar to what they have been trained on
- **Lossy** → decompressed outputs will be degraded compared to the original inputs
- **Learn automatically** → easy to train, no feature engineering required

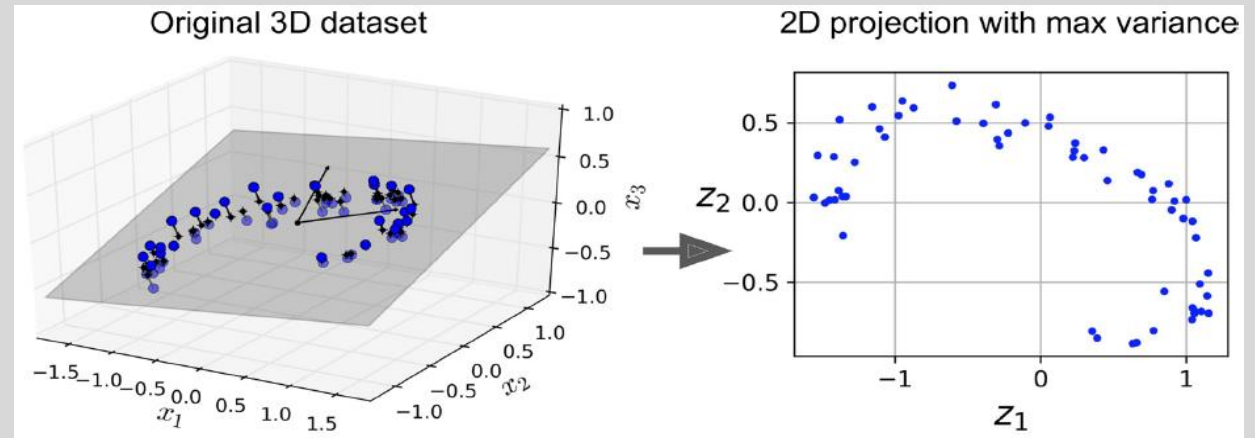


Three important components of an AE model

- **Encoding function** → Neural Network
- **Decoding function** → Neural Network
- **Loss function** → distance between input and generated output

UNDERCOMPLETE LINEAR AUTOENCODER (AE)

- Shallow Machine Learning
- uses only linear activations
- cost function is the mean squared error
- Similar to PCA



```
from tensorflow import keras
```

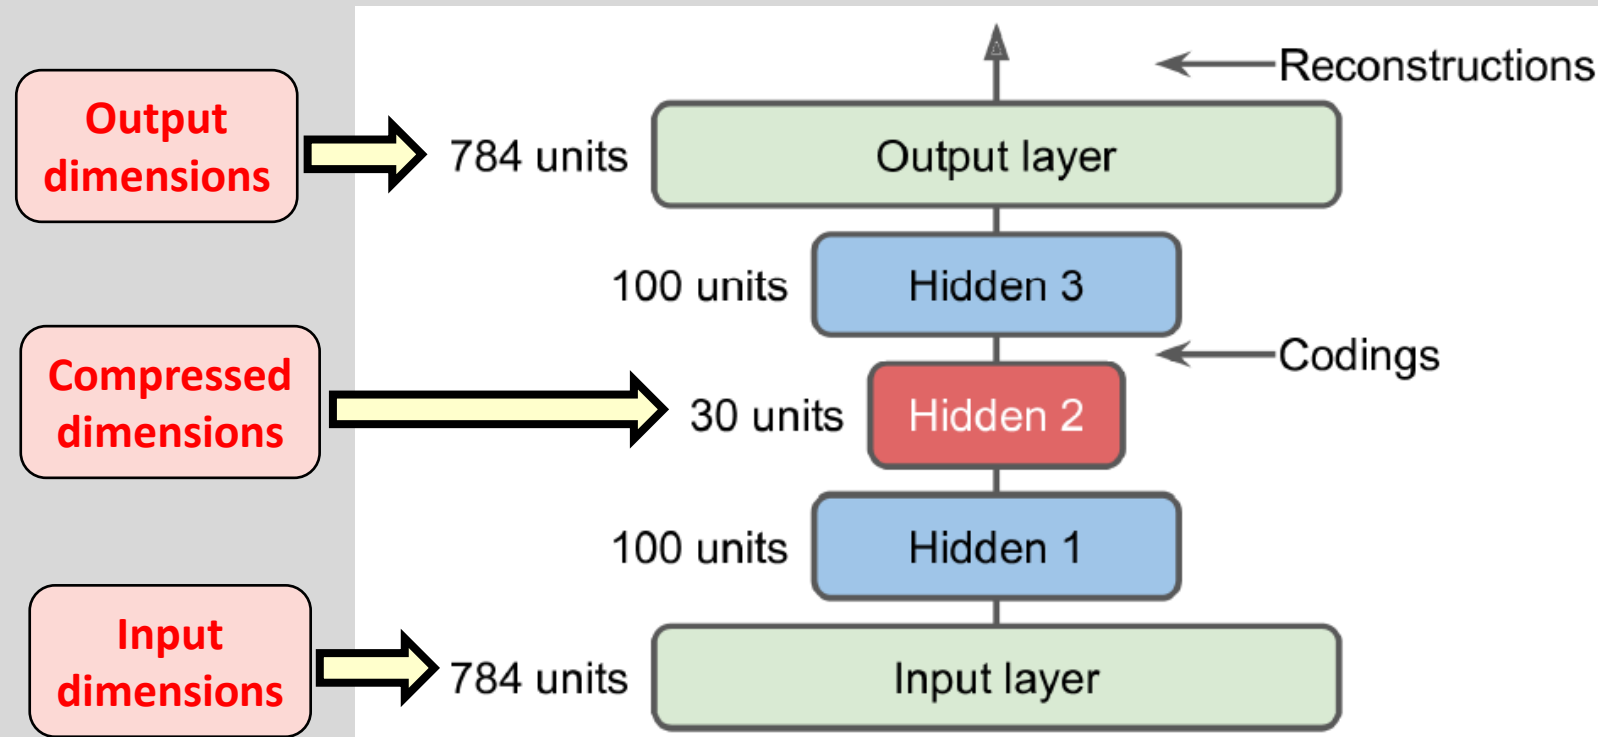
```
encoder = keras.models.Sequential([keras.layers.Dense(2, input_shape=[3])])  
decoder = keras.models.Sequential([keras.layers.Dense(3, input_shape=[2])])  
autoencoder = keras.models.Sequential([encoder, decoder])
```

```
autoencoder.compile(loss="mse", optimizer=keras.optimizers.SGD(lr=0.1))
```

```
history = autoencoder.fit(X_train, X_train, epochs=20)  
codings = encoder.predict(X_train)
```

STACKED/DEEP AUTOENCODER (SAE/DAE)

- have multiple hidden layers
- learn more complex codings
- Symmetric architecture → central hidden layer or coding layer
- sandwich



STACKED/DEEP AUTOENCODER EXAMPLE

- Two Deep Neural Networks → Encoder and Decoder
- Multiple hidden layers for each network

```
stacked_encoder = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    keras.layers.Dense(100, activation="selu"),
    keras.layers.Dense(30, activation="selu"),
])
stacked_decoder = keras.models.Sequential([
    keras.layers.Dense(100, activation="selu", input_shape=[30]),
    keras.layers.Dense(28 * 28, activation="sigmoid"),
    keras.layers.Reshape([28, 28])
])
stacked_ae = keras.models.Sequential([stacked_encoder, stacked_decoder])
stacked_ae.compile(loss="binary_crossentropy",
                   optimizer=keras.optimizers.SGD(lr=1.5))
history = stacked_ae.fit(X_train, X_train, epochs=10,
                        validation_data=[X_valid, X_valid])
```

VISUALIZING THE RECONSTRUCTION OF DAE MODEL

```
def plot_image(image):  
    plt.imshow(image, cmap="binary")  
    plt.axis("off")  
  
def show_reconstructions(model, n_images=5):  
    reconstructions = model.predict(X_valid[:n_images])  
    fig = plt.figure(figsize=(n_images * 1.5, 3))  
    for image_index in range(n_images):  
        plt.subplot(2, n_images, 1 + image_index)  
        plot_image(X_valid[image_index])  
        plt.subplot(2, n_images, 1 + n_images + image_index)  
        plot_image(reconstructions[image_index])  
  
show_reconstructions(stacked_ae)
```



Original images (top) and their reconstructions (bottom)

CONVOLUTIONAL AUTOENCODER

- work with huge image datasets
- unsupervised pretraining or dimensionality reduction
- encoder is a regular CNN composed of convolutional layers and pooling layers
- **Encoder** → reduces the inputs (height and width), increases the depth (feature maps)
- **Decoder** → reverse (transpose) of encoder, upscale image and convert to original dimensions

```
conv_encoder = keras.models.Sequential([
    keras.layers.Reshape([28, 28, 1], input_shape=[28, 28]),
    keras.layers.Conv2D(16, kernel_size=3, padding="same", activation="selu"),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Conv2D(32, kernel_size=3, padding="same", activation="selu"),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Conv2D(64, kernel_size=3, padding="same", activation="selu"),
    keras.layers.MaxPool2D(pool_size=2)
])

conv_decoder = keras.models.Sequential([
    keras.layers.Conv2DTranspose(32, kernel_size=3, strides=2, padding="valid",
                                activation="selu",
                                input_shape=[3, 3, 64]),

    keras.layers.Conv2DTranspose(16, kernel_size=3, strides=2, padding="same",
                                activation="selu"),
    keras.layers.Conv2DTranspose(1, kernel_size=3, strides=2, padding="same",
                                activation="sigmoid"),
    keras.layers.Reshape([28, 28])
])

conv_ae = keras.models.Sequential([conv_encoder, conv_decoder])
```

Simple Conv-AE for Fashion MNIST dataset

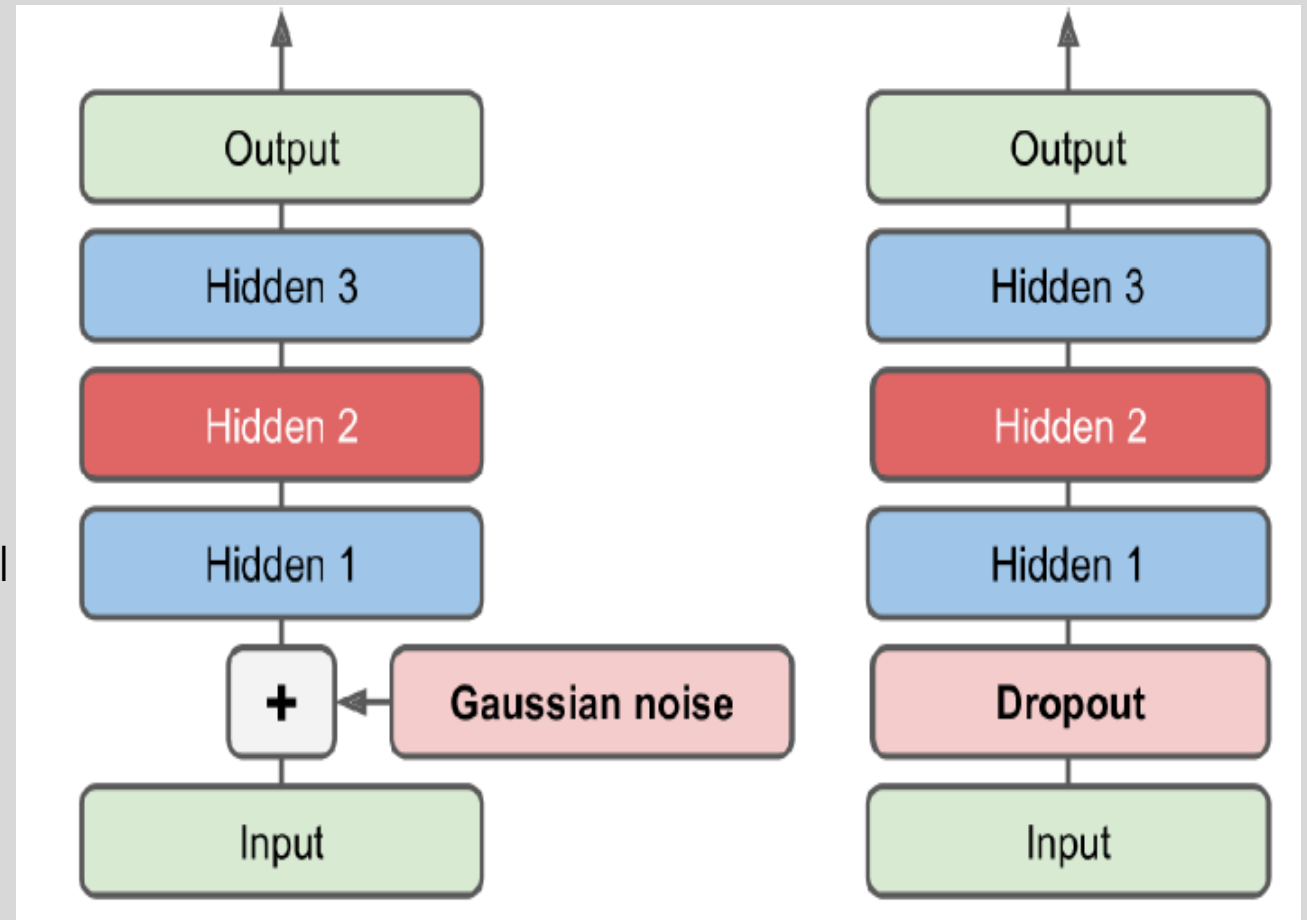
RECURRENT (SEQUENCE-TO-SEQUENCE) AUTOENCODER

- autoencoder for sequences, such as time series or textual information
- unsupervised learning or dimensionality reduction
- **Encoder**→ a sequence-to-vector RNN which compresses the input sequence down to a single vector
- **Decoder**→ vector-to-sequence RNN that does the reverse

```
recurrent_encoder = keras.models.Sequential([
    keras.layers.LSTM(100, return_sequences=True, input_shape=[None, 28]),
    keras.layers.LSTM(30)
])
recurrent_decoder = keras.models.Sequential([
    keras.layers.RepeatVector(28, input_shape=[30]),
    keras.layers.LSTM(100, return_sequences=True),
    keras.layers.TimeDistributed(keras.layers.Dense(28, activation="sigmoid"))
])
recurrent_ae = keras.models.Sequential([recurrent_encoder, recurrent_decoder])
```


DENOISING AUTOENCODER

- learn useful features by adding noise to its inputs
- training it to recover the original, noise-free inputs
- Noise → pure Gaussian Noise or randomly adding dropout
- Regular stacked autoencoder → additional Dropout layer or GaussianNoise layer in encoder



DENOISING AUTOENCODER EXAMPLE

```
dropout_encoder = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(100, activation="selu"),
    keras.layers.Dense(30, activation="selu")
])
dropout_decoder = keras.models.Sequential([
    keras.layers.Dense(100, activation="selu", input_shape=[30]),
    keras.layers.Dense(28 * 28, activation="sigmoid"),
    keras.layers.Reshape([28, 28])
])
dropout_ae = keras.models.Sequential([dropout_encoder, dropout_decoder])
```



Noisy images (top) and their reconstructions (bottom)

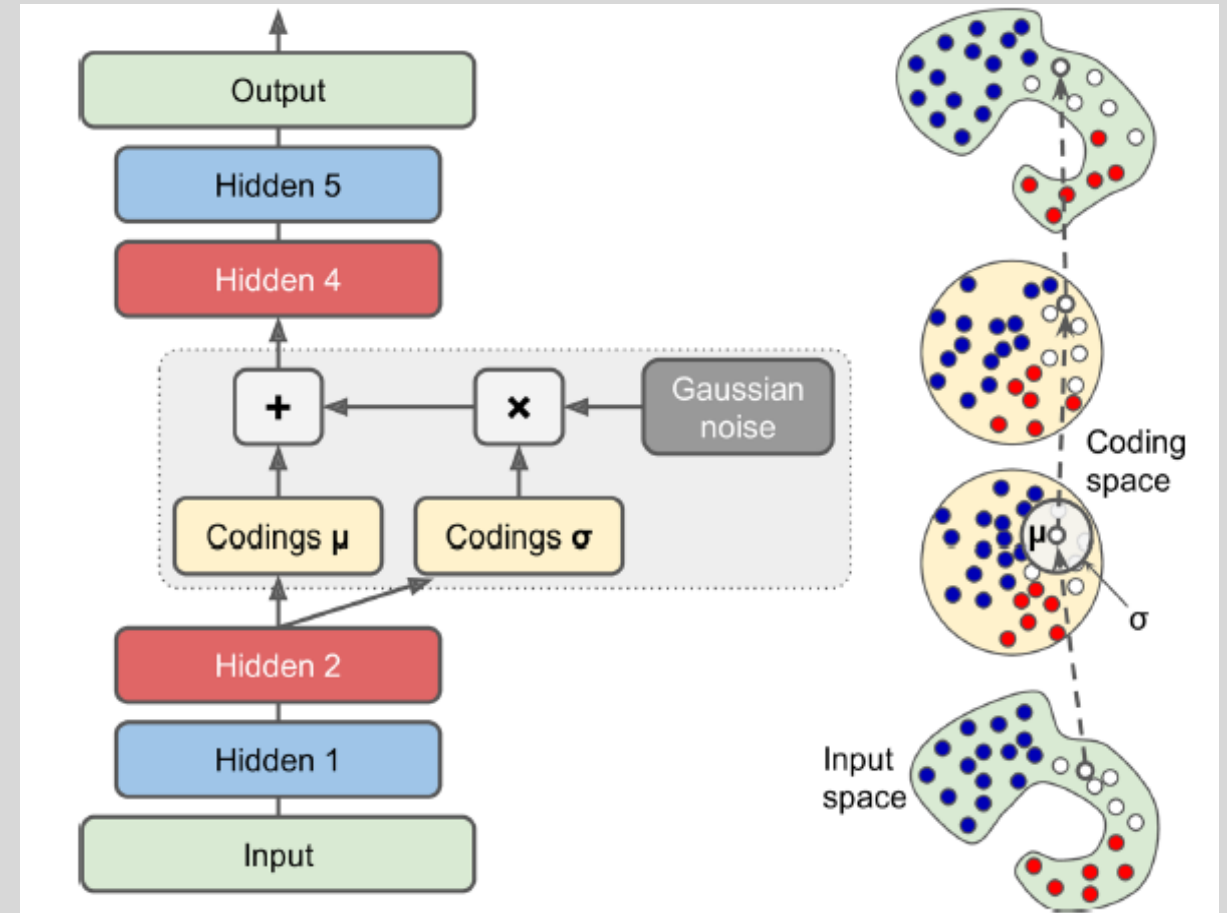
SPARSE AUTOENCODER

- **Sparsity**: adding an appropriate term to the cost function
- reduce the number of active neurons in the coding layer
- For example, AE may be pushed to have on average only 5% significantly active neurons in the coding layer
- Small number of activations
- Each neuron in coding layer → represents a useful feature
- Use **Sigmoid activation** (between 0 and 1) in coding layer
- Add **L1 Regularization** to coding layer's activations
- Decoder → regular decoder

```
sparse_l1_encoder = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    keras.layers.Dense(100, activation="selu"),
    keras.layers.Dense(300, activation="sigmoid"),
    keras.layers.ActivityRegularization(l1=1e-3)
])
sparse_l1_decoder = keras.models.Sequential([
    keras.layers.Dense(100, activation="selu", input_shape=[300]),
    keras.layers.Dense(28 * 28, activation="sigmoid"),
    keras.layers.Reshape([28, 28])
])
sparse_l1_ae = keras.models.Sequential([sparse_l1_encoder, sparse_l1_decoder])
```

VARIATIONAL AUTOENCODER (VAE)

- Introduced in **2013**
- **Diederik Kingma** and **Max Welling**
- most popular generative AE model
- probabilistic autoencoders
- their outputs are partly determined by chance, even after training
- can generate new instances that look like they were sampled from the training set
- instead of directly producing a coding for a given input, the encoder produces a mean coding μ and a standard deviation σ
- Actual coding \rightarrow sampled randomly \rightarrow Gaussian distribution (μ, σ)
- **Latent space/Sampling Layer** \rightarrow cloud of Gaussian points
- very easily generate a new instance
- sample a random coding from the Gaussian distribution and then decode it



Variational autoencoder (left) and an instance going through it (right)

VAE LOSS FUNCTION

Composed of two parts:-

- **Reconstruction Loss:** Loss between the inputs and reconstructions (e.g., *cross entropy*)
- **Latent Loss:** *KL Divergence* between the target distribution (i.e., the Gaussian distribution) and the actual distribution of the codings

The objective function for a VAE model is given below:

$$\log P(X) - D_{KL}[Q(z|X) || P(z|X)] = E[\log P(X|z)] - D_{KL}[Q(z|X) || P(z)]$$

- $\log P(X) \rightarrow$ model our data
- $Q(z | X)$ is the encoder network
- z is the encoded representation
- $P(X | z)$ is the decoder network
- $DKL[Q(z | X) || P(z | X)]$ is the overall error
- $E[\log P(X | z)] \rightarrow$ Maximum Likelihood Estimation \rightarrow given an input z and an output $X \rightarrow$ maximize the conditional distribution $P(X | z)$
- $Q(z | X) \rightarrow$ simple distribution
- $P(z) \rightarrow$ true latent distribution

CONDITIONAL VARIATIONAL AUTOENCODER (CVAE)

- Enhanced VAE model with a conditional variable
- **Kihyuk Sohn, Xincheng Yan, and Honglak Lee** in 2015
- VAE → no control on the data generation process
- problematic if we want to generate some specific data
- For example, generate multiple versions of the handwritten MNIST digit “5”

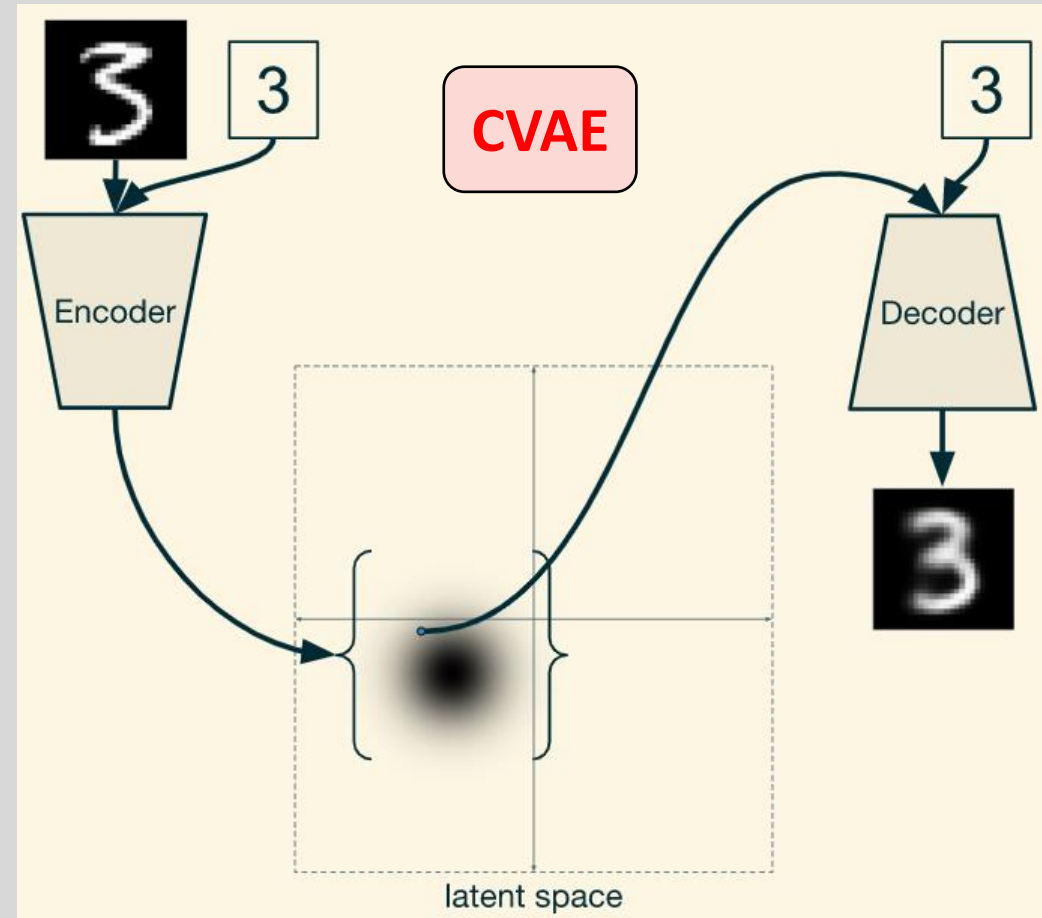
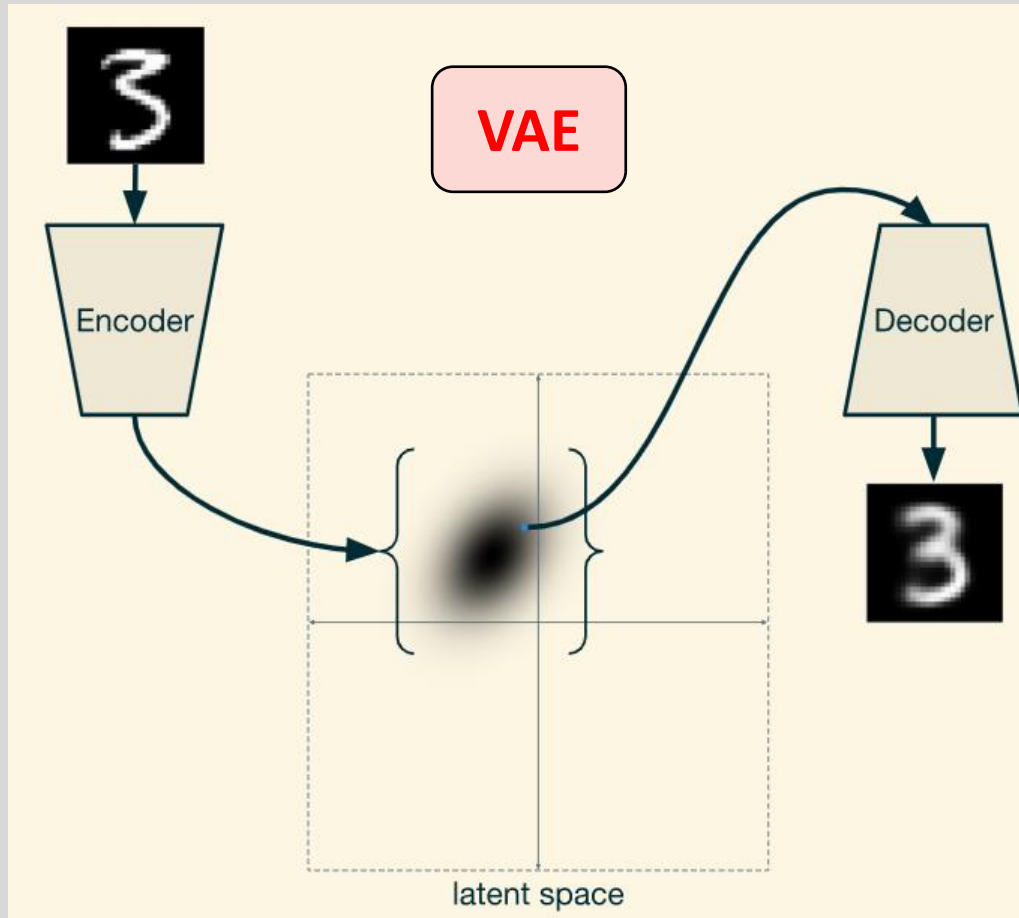


- The **objective function** for a CVAE model is similar to VAE but is conditioned using a variable “c”

$$\log P(X|c) - D_{KL}[Q(z|X, c) \| P(z|X, c)] = E[\log P(X|z, c)] - D_{KL}[Q(z|X, c) \| P(z|c)]$$

- we just conditioned all the distributions with a variable “c”
- Conditional probability distribution (CPD)

VAE VERSUS CVAE



ANOMALY/ OUTLIER DETECTION

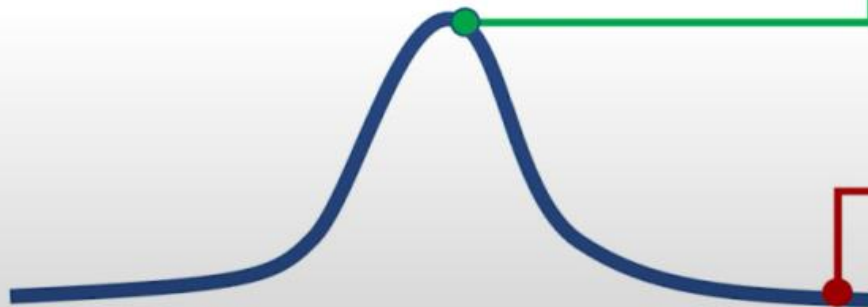
- DAE models → detect rare anomalies or outliers in the data distribution

95% of Driving Data:

(1) sunny, (2) highway, (3) straight road



Detect outliers to avoid unpredictable behavior when training



Edge Cases



Harsh Weather



Pedestrians

VAE PROGRAMMING EXERCISE

(Anomaly Detection)



THANKS!

**Do you have any
questions?**